

Name - Venkatesh Gudikoti

Project: Insider Threat Detection (ML + DL)

Section 1: Context

One-paragraph description


This project focuses on detecting insider threats by analyzing employee behavioral data such as login times, file access frequency, device usage, and access locations. Insider threats are challenging because the users are authorized and malicious behavior often appears subtle. The system learns normal behavior patterns and flags deviations using a combination of machine learning and deep learning-based anomaly detection, producing a risk score and alert for security teams.

Primary technical constraints

- Very limited labeled insider-attack data
- High variation in “normal” behavior across employees
- Need to minimize false positives to avoid alert fatigue
- Privacy-sensitive logs limiting feature granularity

Section 2: Technical Implementation

Architecture Diagram



DL Anomaly Detection (Autoencoder)

↓

Risk Scoring Engine

↓

Threat Alert / Normal Status

Code Walk-through: Critical Function

```
def detect_insider_threat(activity):
    features = preprocess(activity)

    ml_score = isolation_forest.decision_function([features])[0]
    reconstructed = autoencoder.predict([features])
    reconstruction_error = np.mean((features - reconstructed[0]) ** 2)

    risk_score = (0.6 * abs(ml_score)) + (0.4 * reconstruction_error)

    if risk_score > THRESHOLD:
        return {
            "status": "THREAT DETECTED",
            "risk_score": round(risk_score * 100, 2),
            "action": "Alert Security Team"
        }
    else:
        return {
            "status": "NORMAL",
            "risk_score": round(risk_score * 100, 2),
            "action": "No Action Required"
        }
```

Data Flow for a Key Operation

1. Employee performs system activities
2. Logs are captured (time, frequency, resource type)
3. Data is preprocessed and normalized
4. Isolation Forest identifies outliers
5. Autoencoder detects abnormal behavior patterns
6. Risk score is calculated
7. Alert generated if threshold is exceeded

Section 3: Technical Decisions (Core)

Decision 1: Anomaly Detection instead of Supervised Learning

- **Chosen:** Isolation Forest + Autoencoder
- **Alternatives:** Logistic Regression, SVM
- **Trade-offs:**
 - Works without labeled attack data
 - Requires careful threshold tuning
- **Reason:** Insider attacks are rare and poorly labeled in real systems.

Decision 2: Hybrid ML + DL Approach

- **Chosen:** Isolation Forest (ML) + Autoencoder (DL)
- **Alternatives:** Only ML or only DL
- **Trade-offs:**
 - Better detection of both global and subtle anomalies
 - Slightly higher computation cost
- **Reason:** Combining models improved reliability and reduced false alerts.

Scaling Bottleneck & Mitigation

- **Bottleneck:** High-volume activity logs per day
- **Mitigation Strategy:**
 - Session-level feature aggregation
 - Batch inference
 - Periodic model retraining

Section 4: Learning & Iteration (Concise)

One technical mistake

Initially, I used a global behavior baseline, which caused frequent false alerts for power users.

What I learned:

Behavior modeling must be **user-specific**, not one-size-fits-all.

One thing I'd do differently today

- Temporal modeling using LSTM
- Role-based behavior baselines
- Explainability dashboards for security analysts

Final Output Example

Threat Status: DETECTED

Risk Score: 88 / 100

Reasons:

- Unusual login time
- Excessive file access
- New device detected

Action: Security Team Alerted

Why This Project Represents My Strengths

This project highlights my ability to:

- Design ML + DL systems for real-world security problems
- Work with unlabeled and no

3 Technical Decisions

Decision 1

Decision: Use anomaly detection instead of supervised classification

Alternatives Considered:

- Logistic Regression
- Support Vector Machine

Trade-offs:

- Works without labeled data
- Requires careful threshold tuning

Outcome:

- Reduced false negatives and handled unseen threats better

Decision 2

Decision: Use Isolation Forest for ML-based anomaly detection

Alternatives Considered:

- One-Class SVM
- Local Outlier Factor

Trade-offs:

- Scales well to large datasets
- Less effective for highly local anomalies

Outcome:

- Provided stable baseline anomaly detection

Decision 3

Decision: Combine ML with Autoencoder (Deep Learning)

Alternatives Considered:

- Only ML
- Only Deep Learning

Trade-offs:

- Better detection accuracy
- Higher compute cost

Outcome:

- Reduced false positives in real user data

b) Technical Code Walkthrough – Critical Function

Function Name

detect_insider_threat(activity)

Purpose of the Function

This function is responsible for detecting suspicious employee behavior by **combining** machine learning–based anomaly detection with deep learning–based behavior reconstruction. It outputs a risk score and a final decision (Threat / Normal).

Code

```
def detect_insider_threat(activity):
    features = preprocess(activity)

    ml_score = isolation_forest.decision_function([features])[0]
    reconstructed = autoencoder.predict([features])
    reconstruction_error = np.mean((features - reconstructed[0]) ** 2)

    risk_score = (0.6 * abs(ml_score)) + (0.4 * reconstruction_error)

    if risk_score > THRESHOLD:
        return {
            "status": "THREAT DETECTED",
            "risk_score": round(risk_score * 100, 2),
            "action": "Alert Security Team"
        }
    else:
        return {
            "status": "NORMAL",
            "risk_score": round(risk_score * 100, 2),
            "action": "No Action Required"
        }
```

Step-by-Step Walkthrough

1. Input Handling

```
features = preprocess(activity)
```

- The function receives **raw employee activity data** such as login time, file access count, device usage, and location.
- The preprocess step converts this raw data into **numerical features** using normalization and encoding so models can process it.

2. ML-Based Anomaly Detection

```
ml_score = isolation_forest.decision_function([features])[0]
```

- An **Isolation Forest** model checks how unusual the activity is compared to normal employee behavior.
- Lower scores indicate **higher anomaly**.
- This captures **global deviations**, such as unusually high access volume.

3. DL-Based Behavior Reconstruction

```
reconstructed = autoencoder.predict([features])
```

```
reconstruction_error = np.mean((features - reconstructed[0]) ** 2)
```

- A **deep learning autoencoder** tries to reconstruct normal behavior.
- If reconstruction error is high, it indicates **subtle or previously unseen abnormal patterns**.
- This helps detect **slow or stealthy insider threats**.

4. Risk Scoring Logic

```
risk_score = (0.6 * abs(ml_score)) + (0.4 * reconstruction_error)
```

- Both ML and DL outputs are combined into a **single risk score**.
- ML has slightly higher weight to capture obvious anomalies, while DL captures nuanced behavior changes.
- This hybrid approach reduces false positives.

5. Decision & Output

```
if risk_score > THRESHOLD:
```

- If the risk score exceeds a defined threshold:
 - The activity is flagged as a **potential insider threat**
 - An alert is triggered
- Otherwise, the activity is considered normal.